

## ▼ Import Libraries

```
import os
import numpy as np
import matplotlib.pyplot as plt

from pathlib import Path
from collections import Counter

import tensorflow as tf
from tensorflow import keras
from keras import layers
```

## ▼ Loading Data

```
!curl -LO https://github.com/AakashKumarNain/CaptchaCracker/raw/m
!unzip -qq captcha_images_v2.zip
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
0 0	0 0	0 0	0 0	--:--:--	--:--:--	--:--:--	0
100 8863k	100 8863k	0 0	13.5M 0	--:--:--	--:--:--	--:--:--	13.5M

```
# Path to the data directory
data_dir = Path("./captcha_images_v2/")
```

```
# Get list of all the images
images = sorted(list(map(str, list(data_dir.glob("*.png")))))
labels = [img.split(os.path.sep)[-1].split(".png")[0] for img in images]
characters = set(char for label in labels for char in label)
characters = sorted(list(characters))
```

```
print("Number of images found: ", len(images))
print("Number of labels found: ", len(labels))
print("Number of unique characters: ", len(characters))
```

```
print("Characters present: ", characters)

# Batch size for training and validation
batch_size = 16

# Desired image dimensions
img_width = 200
img_height = 50

# Factor by which the image is going to be downsampled
# by the convolutional blocks. We will be using two
# convolution blocks and each block will have
# a pooling layer which downsample the features by a factor of 2.
# Hence total downsampling factor would be 4.
downsample_factor = 4

# Maximum length of any captcha in the dataset
max_length = max([len(label) for label in labels])
```

```
Number of images found: 1040
Number of labels found: 1040
Number of unique characters: 19
Characters present: ['2', '3', '4', '5', '6', '7', '8', 'b', 'c', 'd', 'e', 'f', 'g',
```



## ▼ Preprocessing

```
# Mapping characters to integers
char_to_num = layers.StringLookup(
    vocabulary=list(characters), mask_token=None
)

# Mapping integers back to original characters
num_to_char = layers.StringLookup(
    vocabulary=char_to_num.get_vocabulary(), mask_token=None, inv
```

```
def split_data(images, labels, train_size=0.9, shuffle=True):
    # 1. Get the total size of the dataset
    size = len(images)
    # 2. Make an indices array and shuffle it, if required
    indices = np.arange(size)
    if shuffle:
        np.random.shuffle(indices)
    # 3. Get the size of training samples
    train_samples = int(size * train_size)
    # 4. Split data into training and validation sets
    x_train, y_train = images[indices[:train_samples]], labels[indices[:train_samples]]
    x_valid, y_valid = images[indices[train_samples:]], labels[indices[train_samples:]]
    return x_train, x_valid, y_train, y_valid
```

```
# Splitting data into training and validation sets
x_train, x_valid, y_train, y_valid = split_data(np.array(images),
```

```
def encode_single_sample(img_path, label):
    # 1. Read image
    img = tf.io.read_file(img_path)
    # 2. Decode and convert to grayscale
    img = tf.io.decode_png(img, channels=1)
    # 3. Convert to float32 in [0, 1] range
    img = tf.image.convert_image_dtype(img, tf.float32)
    # 4. Resize to the desired size
    img = tf.image.resize(img, [img_height, img_width])
    # 5. Transpose the image because we want the time
    # dimension to correspond to the width of the image.
    img = tf.transpose(img, perm=[1, 0, 2])
    # 6. Map the characters in label to numbers
    label = char_to_num(tf.strings.unicode_split(label, input_enc))
    # 7. Return a dict as our model is expecting two inputs
    return {"image": img, "label": label}
```

## ▼ Create Dataset Objects

```
train_dataset = tf.data.Dataset.from_tensor_slices((x_train, y_train))
train_dataset = (
    train_dataset.map(
        encode_single_sample, num_parallel_calls=tf.data.AUTOTUNE
    )
    .batch(batch_size)
    .prefetch(buffer_size=tf.data.AUTOTUNE)
)


validation_dataset = tf.data.Dataset.from_tensor_slices((x_valid, y_valid))
validation_dataset = (
    validation_dataset.map(
        encode_single_sample, num_parallel_calls=tf.data.AUTOTUNE
    )
    .batch(batch_size)
    .prefetch(buffer_size=tf.data.AUTOTUNE)
)
```

## ▼ Visualization

```
_, ax = plt.subplots(4, 4, figsize=(12, 8))
for batch in train_dataset.take(1):
    images = batch["image"]
    labels = batch["label"]
    for i in range(16):
        img = (images[i] * 255).numpy().astype("uint8")
        label = tf.strings.reduce_join(num_to_char(labels[i])).numpy()
        ax[i // 4, i % 4].imshow(img[:, :, 0].T, cmap="gray")
        ax[i // 4, i % 4].set_title(label)
        ax[i // 4, i % 4].axis("off")
plt.show()
```



## ▼ Model



```
class CTCLayer(layers.Layer):
    def __init__(self, name=None):
        super().__init__(name=name)
        self.loss_fn = keras.backend.ctc_batch_cost

    def call(self, y_true, y_pred):
        # Compute the training-time loss value and add it
        # to the layer using `self.add_loss()`.
        batch_len = tf.cast(tf.shape(y_true)[0], dtype="int64")
        input_length = tf.cast(tf.shape(y_pred)[1], dtype="int64")
        label_length = tf.cast(tf.shape(y_true)[1], dtype="int64")

        input_length = input_length * tf.ones(shape=(batch_len, 1), dtype="int64")
        label_length = label_length * tf.ones(shape=(batch_len, 1), dtype="int64")

        loss = self.loss_fn(y_true, y_pred, input_length, label_length)
        self.add_loss(loss)

        # At test time, just return the computed predictions
        return y_pred
```

```
def build_model():
```

```
# Inputs to the model
input_img = layers.Input(
    shape=(img_width, img_height, 1), name="image", dtype="float32")
labels = layers.Input(name="label", shape=(None,), dtype="float32")

# First conv block
x = layers.Conv2D(
    32,
    (3, 3),
    activation="relu",
    kernel_initializer="he_normal",
    padding="same",
    name="Conv1",
)(input_img)
x = layers.MaxPooling2D((2, 2), name="pool1")(x)

# Second conv block
x = layers.Conv2D(
    64,
    (3, 3),
    activation="relu",
    kernel_initializer="he_normal",
    padding="same",
    name="Conv2",
)(x)
x = layers.MaxPooling2D((2, 2), name="pool2")(x)

# We have used two max pool with pool size and strides 2.
# Hence, downsampled feature maps are 4x smaller. The number
# filters in the last layer is 64. Reshape accordingly before
# passing the output to the RNN part of the model
new_shape = ((img_width // 4), (img_height // 4) * 64)
x = layers.Reshape(target_shape=new_shape, name="reshape")(x)
x = layers.Dense(64, activation="relu", name="dense1")(x)
x = layers.Dropout(0.2)(x)

# RNNs
x = layers.Bidirectional(layers.LSTM(128, return_sequences=True
```

```

x = layers.Bidirectional(layers.LSTM(64, return_sequences=True

# Output layer
x = layers.Dense(
    len(char_to_num.get_vocabulary()) + 1, activation="softma
)(x)

# Add CTC layer for calculating CTC loss at each step
output = CTCLayer(name="ctc_loss")(labels, x)

# Define the model
model = keras.models.Model(
    inputs=[input_img, labels], outputs=output, name="ocr_mod
)
# Optimizer
opt = keras.optimizers.Adam()
# Compile the model and return
model.compile(optimizer=opt)
return model

```

```

# Get the model
model = build_model()
model.summary()

```

Model: "ocr\_model\_v1"

Layer (type)	Output Shape	Param #	Connected to
image (InputLayer)	[(None, 200, 50, 1)]	0	[]
Conv1 (Conv2D)	(None, 200, 50, 32)	320	['image[0][0]']
pool1 (MaxPooling2D)	(None, 100, 25, 32)	0	['Conv1[0][0]']
Conv2 (Conv2D)	(None, 100, 25, 64)	18496	['pool1[0][0]']
pool2 (MaxPooling2D)	(None, 50, 12, 64)	0	['Conv2[0][0]']
reshape (Reshape)	(None, 50, 768)	0	['pool2[0][0]']
dense1 (Dense)	(None, 50, 64)	49216	['reshape[0][0]']
dropout_1 (Dropout)	(None, 50, 64)	0	['dense1[0][0]']

```

bidirectional_2 (Bidirectional (None, 50, 256) 197632 ['dropout_1[0][0]']
)

bidirectional_3 (Bidirectional (None, 50, 128) 164352 ['bidirectional_2[0][0]']
)

label (InputLayer) [(None, None)] 0 []

dense2 (Dense) (None, 50, 21) 2709 ['bidirectional_3[0][0]']

ctc_loss (CTCLayer) (None, 50, 21) 0 ['label[0][0]',
'dense2[0][0]']

=====
Total params: 432,725
Trainable params: 432,725
Non-trainable params: 0

```



## ▼ Training

```

epochs = 120
early_stopping_patience = 8
# Add early stopping
early_stopping = keras.callbacks.EarlyStopping(
    monitor="val_loss", patience=early_stopping_patience, restore
)

# Train the model
history = model.fit(
    train_dataset,
    validation_data=validation_dataset,
    epochs=epochs,
    callbacks=[early_stopping],
)

59/59 [=====] - 18s 303ms/step - loss: 0.0170 - val_loss: 0.0170
Epoch 58/120
59/59 [=====] - 19s 316ms/step - loss: 0.0176 - val_loss: 0.0176
Epoch 59/120
59/59 [=====] - 18s 300ms/step - loss: 0.0157 - val_loss: 0.0157
Epoch 60/120
59/59 [=====] - 17s 292ms/step - loss: 0.0132 - val_loss: 0.0132
Epoch 61/120
59/59 [=====] - 17s 292ms/step - loss: 0.0396 - val_loss: 0.0396

```



```
Epoch 62/120
59/59 [=====] - 17s 296ms/step - loss: 0.0239 - val_loss: 0.
Epoch 63/120
59/59 [=====] - 17s 291ms/step - loss: 0.0116 - val_loss: 0.
Epoch 64/120
59/59 [=====] - 17s 293ms/step - loss: 0.0338 - val_loss: 0.
Epoch 65/120
59/59 [=====] - 17s 293ms/step - loss: 0.0440 - val_loss: 0.
Epoch 66/120
59/59 [=====] - 19s 314ms/step - loss: 0.0308 - val_loss: 0.
Epoch 67/120
59/59 [=====] - 17s 294ms/step - loss: 0.0360 - val_loss: 0.
Epoch 68/120
59/59 [=====] - 17s 291ms/step - loss: 0.0246 - val_loss: 0.
Epoch 69/120
59/59 [=====] - 17s 289ms/step - loss: 0.0290 - val_loss: 0.
Epoch 70/120
59/59 [=====] - 17s 291ms/step - loss: 0.0348 - val_loss: 0.
Epoch 71/120
59/59 [=====] - 17s 290ms/step - loss: 0.0883 - val_loss: 0.
Epoch 72/120
59/59 [=====] - 17s 288ms/step - loss: 0.0758 - val_loss: 0.
Epoch 73/120
59/59 [=====] - 17s 290ms/step - loss: 0.0207 - val_loss: 0.
Epoch 74/120
59/59 [=====] - 18s 312ms/step - loss: 0.0141 - val_loss: 0.
Epoch 75/120
59/59 [=====] - 17s 290ms/step - loss: 0.0258 - val_loss: 0.
Epoch 76/120
59/59 [=====] - 17s 291ms/step - loss: 0.0257 - val_loss: 0.
Epoch 77/120
59/59 [=====] - 17s 289ms/step - loss: 0.0463 - val_loss: 0.
Epoch 78/120
59/59 [=====] - 17s 291ms/step - loss: 0.0564 - val_loss: 0.
Epoch 79/120
59/59 [=====] - 17s 291ms/step - loss: 0.0225 - val_loss: 0.
Epoch 80/120
59/59 [=====] - 17s 293ms/step - loss: 0.0608 - val_loss: 0.
Epoch 81/120
59/59 [=====] - 17s 287ms/step - loss: 0.0897 - val_loss: 0.
Epoch 82/120
59/59 [=====] - 18s 311ms/step - loss: 0.0291 - val_loss: 0.
Epoch 83/120
59/59 [=====] - 17s 287ms/step - loss: 0.0267 - val_loss: 0.
Epoch 84/120
59/59 [=====] - 17s 290ms/step - loss: 0.0240 - val_loss: 0.
Epoch 85/120
59/59 [=====] - 17s 287ms/step - loss: 0.0215 - val_loss: 0.▼
```

## ▼ Inference

# Get the prediction model by extracting layers till the output 1

```
prediction_model = keras.models.Model(
    model.get_layer(name="image").input, model.get_layer(name="de
)
prediction_model.summary()
```

# A utility function to decode the output of the network

```
def decode_batch_predictions(pred):
    input_len = np.ones(pred.shape[0]) * pred.shape[1]
    # Use greedy search. For complex tasks, you can use beam search
    results = keras.backend.ctc_decode(pred, input_length=input_len
        :, :max_length
    ]
    # Iterate over the results and get back the text
    output_text = []
    for res in results:
        res = tf.strings.reduce_join(num_to_char(res)).numpy().de
        output_text.append(res)
    return output_text
```

# Let's check results on some validation samples

```
for batch in validation_dataset.take(1):
    batch_images = batch["image"]
    batch_labels = batch["label"]

    preds = prediction_model.predict(batch_images)
    pred_texts = decode_batch_predictions(preds)

    orig_texts = []
    for label in batch_labels:
        label = tf.strings.reduce_join(num_to_char(label)).numpy()
        orig_texts.append(label)

_, ax = plt.subplots(4, 4, figsize=(15, 5))
for i in range(len(pred_texts)):
    img = (batch_images[i, :, :, 0] * 255).numpy().astype(np.
    img = img.T
    title = f"Prediction: {pred_texts[i]}"
    ax[i // 4, i % 4].imshow(img, cmap="gray")
    ax[i // 4, i % 4].set_title(title)
```

```
ax[i // 4, i % 4].set_title(title)
ax[i // 4, i % 4].axis("off")
plt.show()
```

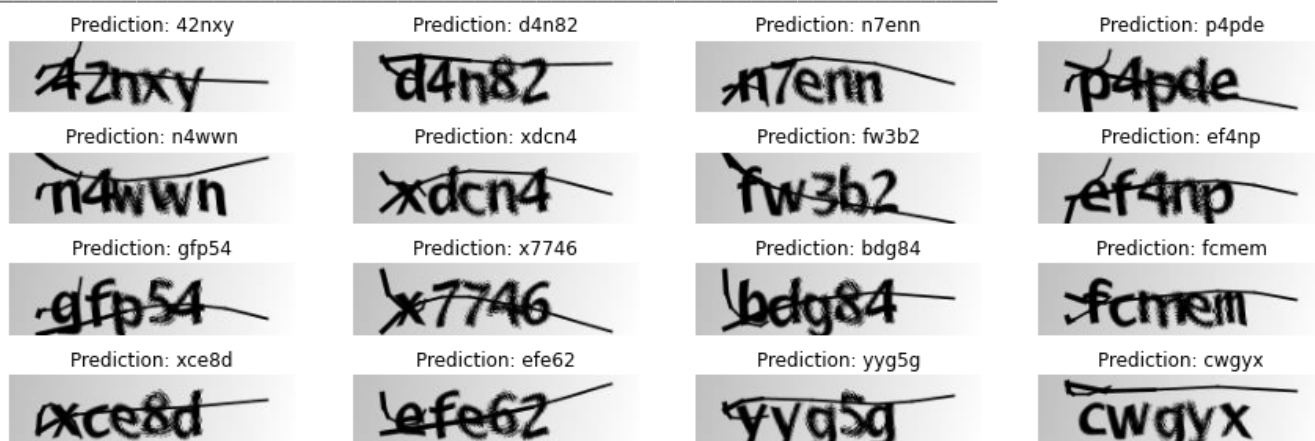
Model: "model\_1"

Layer (type)	Output Shape	Param #
image (InputLayer)	[(None, 200, 50, 1)]	0
Conv1 (Conv2D)	(None, 200, 50, 32)	320
pool1 (MaxPooling2D)	(None, 100, 25, 32)	0
Conv2 (Conv2D)	(None, 100, 25, 64)	18496
pool2 (MaxPooling2D)	(None, 50, 12, 64)	0
reshape (Reshape)	(None, 50, 768)	0
dense1 (Dense)	(None, 50, 64)	49216
dropout_1 (Dropout)	(None, 50, 64)	0
bidirectional_2 (Bidirectional)	(None, 50, 256)	197632
bidirectional_3 (Bidirectional)	(None, 50, 128)	164352
dense2 (Dense)	(None, 50, 21)	2709

=====  
Total params: 432,725

Trainable params: 432,725

Non-trainable params: 0



Colab paid products - Cancel contracts here

✓ 3s completed at 14:49

